

BFGS Algorithm for finding minimizer

Line search procedure

The method is adapted directly from *Numerical Optimization* by Jorge Nocedal and Stephen J. Wright, called the *strong backtracking*. The procedure will be explained later.

1. Set $a = 1$. Set $a_{old} = 0$, then go to 2.
2. If $f(x + as) > f(x) + \mu as^T \nabla f(x)$ or $f(x + as) > f(x + a_{old}s)$, then set $a_{low} = a_{old}$ and $a_{high} = a$ and go to step 6. Else, go to step 3.
3. If $|s^T \nabla f(x + as)| \leq -\eta s^T \nabla f(x)$, then $\lambda = a$ and Exit. Else, go to step 4.
4. If $s^T \nabla f(x + as) \geq 0$ then set $a_{low} = a_{old}$ and $a_{high} = a$ and go to step 6. Else, go to step 5.
5. Let $a_{old} = a$ and $a = 2a$. Back to step 2.
6. Set $f_{low} = f(x + a_{low}s)$. Go to step 7. Note that this value will be fixed regardless of changing a_{low} .
7. Use binary search or golden section search to find a suitable a between a_{low} and a_{high} .
 - (a) For binary search, let $a = (a_{low} + a_{high})/2$
 - (b) For golden section search, let $c = a_{low} + (a_{high} - a_{low})/\phi$ and $d = a_{high} - (a_{high} - a_{low})/\phi$, where ϕ is a golden ratio. If $f(x + cs) < f(x + ds)$, let $a = c$. Otherwise, let $a = d$.

Then go to step 8.

8. If $f(x + as) > f(x) + \mu as^T \nabla f(x)$ or $f(x + as) > f_{low}$, then let $a_{high} = a$. Then go back to step 7. Else, go to step 9.
9. If $|s^T \nabla f(x + as)| \leq -\eta s^T \nabla f(x)$, then $\lambda = a$ and Exit. Else, go to step 10.
10. If $(s^T \nabla f(x + as))(a_{high} - a_{low}) \geq 0$, then let $a_{high} = a_{low}$ and goto step 11.
11. Let $a_{low} = a$ and back to step 7.

The step-by-step explanation is as follows:

1. a is representing λ . Let it be 1 first.
2. If $f(x + as) > f(x) + \mu as^T \nabla f(x)$, then it is violating first Wolfe's condition. It is thus ensuring that the bracket (a_{old}, a) will contain a range that does not violate that. (The slope at $s^T \nabla f(x + a_{old}s)$ is always negative according to step 4, so that (a_{old}, a) must contain local minimum, that is, containing second Wolfe's point.) Also, $f(x + as) > f(x + a_{old}s)$ is indicating that the function is going to increase, thus (a_{old}, a) must also contain local minimum (that is second Wolfe's point, too.) Note that a_{high} and a_{low} can swap regardless of their values.
3. The first Wolfe's condition is already checked in step 2. Thus, if the point also satisfies the second Wolfe's condition, then let it be λ and exit.
4. If $s^T \nabla f(x + as) \geq 0$ then it is indicating that the function is going to increase, thus (a_{old}, a) must contain local minimum (that is second Wolfe's point.) (It is cleared because $s^T \nabla f(x + a_{old}s)$ is always negative.) Also, note that a_{high} and a_{low} can swap regardless of their values.
5. If the range (a_{old}, a) does not contain the second Wolfe's point, slide it to the immediate right and expand it two times.
6. This evaluated value will be used in step 8.
7. Finding a between a_{high} and a_{low} . a will converge to a range that satisfies the second Wolfe's condition.

8. If $f(x + as) > f(x) + \mu as^T \nabla f(x)$ or $f(x + as) > f_{low}$, then $f(x + as)$ value is too big. Thus, let a_{high} be it to lower the upper bound.
9. The first Wolfe's condition is already checked in step 8. Thus, if the point also satisfies the second Wolfe's condition, then let it be λ and exit.
10. $(s^T \nabla f(x + as))(a_{high} - a_{low}) \geq 0$, together with step 11, ensures us that the range between a_{high} and a_{low} always contains local minimum(s). To see that, let $a_{high} < a_{low}$, so $a_{high} - a_{low} < 0$. Moreover, the slope at a_{high} must be negative and the slope at a_{low} must be positive (from step 2, step 4, and recursive characteristics that happen here.) $a_{high} < a < a_{low}$ must be true, and in step 11 a_{low} becomes a , so the slope at a must be positive. If the slope at a is negative, then a_{high} must be swapped with a_{low} to make the slope of a_{high} positive. Then between a_{high} and a will be ensured to have a local minimum. The same is applied when $a_{high} > a_{low}$. Just swap the pair accordingly, and we will get the same proof.
11. Let $a_{low} = a$ to shorten the length.

From the book *Numerical Optimization*, steps 2 to 5 are called *Bracket phase*, and steps 7 to 11 are called *Zoom phase*.

BFGS algorithm

First is an implementation of `StrongBacktrack.m`. The explanation is already given in greater detail in the last section. I chose to use the binary search because it usually takes fewer steps than the golden section search.

```

1 function [lambda,nF,nG] = StrongBacktrack(FcnName, x0, s, a, mu, eta)
2
3 % ----- Function inputs -----
4 %
5 % FcnName: function to return the value, the gradient, and the Hessian
6 % of the particular function.
7 %   Mode 1: return only f.
8 %   Mode 2: return f and gradient.
9 %
10 % x0: starting point of searching.
11 %
12 % s: search direction.
13 %
14 % a: size of initial lambda.
15 %
16 % mu, eta: the parameters used in the stopping criterion for line search.
17
18 % ----- Function outputs -----
19 %
20 % lambda: returned the lambda value that satisfies strong Wolfe's.
21 %
22 % nF, nG: numbers of f and gradient calculations.
23
24 nF = 0; % number of f calculations.
25 nG = 0; % number of gradient calculations.
26
27 [f0, g0] = FcnName(x0, 2); nF = nF + 1; nG = nG + 1;
28 % first evaluation of f and gradient.
29 fprev = f0; aprev = 0;
30 % fprev: previous value of f at (x0 + a*s).
31 % aprev: previous value of a (that is the value going to represent lambda.)
32 alo = NaN; ahi = NaN; % the bracket that contains strong Wolfe's.
33 % alo represent lower value of f, ahi represent higher value of f.
34 % However, the value of alo and ahi can be swapped without ruining the algorithm.
35
36 gr = (sqrt(5) + 1) / 2; % golden ratio.
37
38 % finding suitable bracket.
```

```

39 while 1
40     [f1, g1] = FcnName(x0 + a*s, 2); nF = nF + 1; nG = nG + 1;
41     % evaluate function at (x0 + a*s).
42     if f1 > f0 + mu*a*dot(s,g0) || f1 >= fprev
43         % If the function value of the right point is more than Armijo's rule or
44         % more than that of previous value,
45         % it is sure that the lowest point is there in the bracket, and
46         % there is strong Wolfe's point in that bracket. (Since it is ensured that
47         % at aprev, the slope is negative.)
48         alo = aprev; ahi = a;
49         break
50     elseif abs(dot(g1,s)) <= -eta*dot(g0,s)
51         % Checking the second strong Wolfe's condition.
52         lambda = a;
53         return
54     elseif dot(g1,s) >= 0
55         % If the slope on the right point is positive, it is sure that the
56         % lowest point is there in the bracket, and
57         % there is strong Wolfe's point in that bracket.
58         % (Since it is ensured that at aprev, the slope is negative.)
59         alo = aprev; ahi = a;
60         break
61     end
62     % slide the bracket to the right and expand it by multiple of 2.
63     fprev = f1; aprev = a; a = 2*a;
64 end
65 [flo,~] = FcnName(x0 + alo*s, 1); nF = nF + 1; % f value at (x0 + alo*s).
67 while 1
68
69     % if want to use golden search, use these lines.
70     % ----Golden search to find a between alo and ahi.
71     % ----f1 is f at (x0 + a*s), g1 is gradient at (x0 + a*s).
72     %     c = ahi - (ahi - alo) / gr;
73     %     d = alo + (ahi - alo) / gr;
74     %
75     %     [fc,gc] = FcnName(x0 + c*s, 2);
76     %     [fd,gd] = FcnName(x0 + d*s, 2);
77     %
78     %     if fc < fd % f(c) > f(d) to find the maximum.
79     %         f1 = fc; g1 = gc; a = c;
80     %     else
81     %         f1 = fd; g1 = gd; a = d;
82     %     end
83     % ----end of Golden search
84
85     % if want to use binary search, use this line.
86     a = (ahi+alo)/2; [f1,g1] = FcnName(x0 + a*s, 2);
87     nF = nF + 1; nG = nG + 1;
88
89     if f1 > f0 + mu*a*dot(s,g0) || f1 > flo
90         % if violating Armijo's rule or if f1 still higher than flo, then,
91         % set new hi to decrease f at ahi.
92         ahi = a;
93     else
94         if abs(dot(g1,s)) <= -eta*dot(g0,s)
95             % Checking the second strong Wolfe's condition.
96             lambda = a;
97             return
98         elseif dot(g1,s)*(ahi - alo) >= 0
99             % This condition ensures that ahi and alo always bracket the
100             % lowest point. That is, there is strong Wolfe's point between
101             % alo and ahi.
102             ahi = alo;
103         end
104         alo = a;
105     end
106

```

```

107 end
108
109 end

```

Implementation of function `BFGS.m` is given here. The `epsilon` value is used to give a threshold for the norm of the gradient (the gradient of the local minimum must converge to zero).

```

1 function [xmin,fmin,Xk,Fk,Gk,Lk,nF,nG,IFLAG] = BFGS(FcnName,x0,epsilon,mu,eta,itmax)
2
3 % ----- Function inputs -----
4 %
5 % FcnName: function to return the value, the gradient, and the Hessian
6 % of the particular function.
7 % Mode 1: return only f.
8 % Mode 2: return f and gradient.
9 %
10 % x0: starting point of searching.
11 %
12 % epsilon: stopping criterion of the minimum search. (norm(x1-x0) < epsilon.)
13 %
14 % mu, eta: the parameters used in the stopping criterion for line search.
15 %
16 % itmax: max allowed number of iterations.
17 %
18 % IFLAG: success (0) or not success (-999).
19
20 % ----- Function outputs -----
21 %
22 % xmin, fmin: returned minimum function argument and value, respectively.
23 %
24 % Xk ,Fk, Gk, Lk: arrays to keep x, f, gradient and lambda along the search steps.
25 %
26 % nF, nG: numbers of f and gradient calculations.
27 %
28 % IFLAG: indicate the success. 0 if success, -999 otherwise.
29
30 Xk = []; % list to store x_k.
31 Fk = []; % list to store f_k.
32 Gk = []; % list to store g_k.
33 Lk = []; % list to store l_k.
34
35 nF = 0; % number of f calculations.
36 nG = 0; % number of gradient calculations.
37
38 IFLAG = -999; % IFLAG: indicate the success.
39
40 B = eye(2); % Let the first matrix B be an identity matrix.
41
42 for i = 1:itmax
43
44     % strong backtracking.
45     [f0, g0] = FcnName(x0, 2); nF = nF + 1; nG = nG + 1;
46     a = 1; % first value of lambda.
47     s = B\(-g0); % set line search direction.
48     [lambda,nFnew,nGnew] = StrongBacktrack(FcnName, x0, s, a, mu, eta);
49     % finding lambda that satisfied strong Wolfe's.
50     nF = nF + nFnew; nG = nG + nGnew;
51
52     % store values.
53     Xk(:,i) = x0; Fk(i) = f0; Gk(:,i) = g0; Lk(i) = lambda;
54
55     % update B.
56     x1 = x0 + lambda*s;
57     [f1,g1] = FcnName(x1, 2); nF = nF + 1; nG = nG + 1;
58     delta_g = g1 - g0;
59     delta_x = lambda*s;
60     B = B + delta_g*delta_g'/dot(delta_g,delta_x) - B*(delta_x*delta_x')*B/(delta_x'*B*
        delta_x);
61

```

```

62 % terminate
63 if norm(g1) < epsilon % at local minimum, gradient converges to 0.
64     xmin = x1; fmin = f1; IFLAG = 0;
65     disp('search successful. ');
66     break
67 end
68
69 % update values
70 x0 = x1; f0 = f1; g0 = g1;
71 end
72
73 if IFLAG == -999
74     xmin = 0; fmin = 0; disp('search unsuccessful. ');
75 end
76
77 end

```

Implementation of Rosenbrock.m is here.

```

1 function [f,gradient] = Rosenbrock(x,options)
2
3 % Declare the functions.
4 f_fun = @(x) 100*(x(2) - x(1)^2)^2 + (1-x(1))^2;
5 gradient_fun = @(x)[400*x(1)*(x(1)^2 - x(2)) - 2*(1-x(1)); 200*(x(2)-x(1)^2)];
6
7 % Evaluate numerical values.
8 switch options
9     case 1 % calculate only f.
10         f = f_fun(x);
11         gradient = 0;
12     case 2 % calculate f and gradient.
13         f = f_fun(x);
14         gradient = gradient_fun(x);
15     otherwise % invalid option.
16         disp('invalid option. ')
17         f = 0; gradient = 0;
18 end
19
20 end

```

This is a script used to test the code.

```

1 [xmin,fmin,Xk,Fk,Gk,Lk,nF,nG,IFLAG] = BFGS(@Rosenbrock,[10;12],0.000002,1e-4,0.95,10000);
2
3 % print out the result.
4 fprintf('%5s %13s %13s %15s %15s %15s \n', 'Iter', 'x_1', 'x_2', 'f', 'gradient_1', 'gradient_2');
5 for i = 0:length(Xk)-1
6     fprintf('%5.2d %13.7f %13.7f %15.5f %15.5f %15.5f \n', i, Xk(1,i+1), Xk(2,i+1), Fk(i+1), Gk(1,i+1), Gk(2,i+1));
7 end
8
9 fprintf("Number of f calculations: %i \n", nF)
10 fprintf("Number of gradient calculations: %i \n", nG)

```

The reported tabular is given here when setting $x_0 = [10;12]$, $\epsilon = 2e-6$, $\mu = 1e-4$, $\eta = 0.1$.

```

1 search successful.
2 Iter      x_1      x_2      f      gradient_1      gradient_2
3 00      10.0000000      12.0000000      774481.00000      352018.00000      -17600.00000
4 01      -0.7427368      12.5371094      14368.14165      3557.32893      2397.09028
5 02      -1.3243799      2.4088612      48.28939      342.27468      130.97580
6 03      -1.3592801      1.8692915      5.61307      7.05237      4.32984
7 04      -1.3603042      1.8516118      5.57118      -4.07623      0.23685
8 05      -1.1513687      1.2774932      4.86029      -26.48116      -9.63133
9 06      -0.9980416      0.9288701      4.44398      -30.83025      -13.44341
10 07      -0.7386240      0.6116981      3.46017      16.06163      13.22654
11 08      -0.5535744      0.2865384      2.45322      -7.51499      -3.98125
12 09      -0.5106769      0.2249715      2.41045      -10.33820      -7.16387
13 10      -0.3238515      0.0558135      1.99333      -9.00379      -9.81327

```

14	11	-0.2558141	0.0796349	1.59722	-1.05922	2.83880
15	12	-0.1089611	-0.0123586	1.28851	-3.27402	-4.84622
16	13	-0.0343320	-0.0372802	1.21775	-2.59681	-7.69178
17	14	0.2377167	0.0245037	0.68351	1.51873	-6.40111
18	15	0.2279841	0.0512564	0.59606	-1.47834	-0.14406
19	16	0.3709929	0.1168103	0.43902	1.83241	-4.16508
20	17	0.4885548	0.2088254	0.35074	4.81248	-5.97208
21	18	0.5326485	0.2882858	0.22051	-1.90869	0.91428
22	19	0.6365177	0.3911262	0.15180	2.84484	-2.80573
23	20	0.7350614	0.5217170	0.10478	4.93849	-3.71967
24	21	0.7609271	0.5815879	0.05782	-1.26277	0.51557
25	22	0.8381920	0.6951434	0.03169	2.16497	-1.48449
26	23	0.8999793	0.7999008	0.02013	3.42218	-2.01239
27	24	0.9190669	0.8464102	0.00685	-0.79648	0.34525
28	25	0.9634323	0.9253881	0.00213	1.01122	-0.56276
29	26	0.9967803	0.9931595	0.00003	0.15760	-0.08228
30	27	0.9967820	0.9935801	0.00001	-0.00868	0.00113
31	28	1.0000038	0.9999972	0.00000	0.00416	-0.00208
32	29	0.9999966	0.9999933	0.00000	-0.00001	0.00000
33	Number of f calculations:			245		
34	Number of gradient calculations:			222		

The reported tabular is given here when setting $x_0 = [10;12]$, $\epsilon = 2e-6$, $\mu = 1e-4$, $\eta = 0.95$.

1	search successful.					
2	Iter	x_1	x_2	f	gradient_1	gradient_2
3	00	10.0000000	12.0000000	774481.00000	352018.00000	-17600.00000
4	01	-0.7427368	12.5371094	14368.14165	3557.32893	2397.09028
5	02	-1.3243799	2.4088612	48.28939	342.27468	130.97580
6	03	-1.3642658	1.7922101	6.06601	-42.38830	-13.80221
7	04	-1.3602443	1.8519908	5.57105	-3.78127	0.34524
8	05	-1.3601578	1.8513708	5.57052	-3.99046	0.26830
9	06	-1.3566856	1.8352821	5.55679	-7.59695	-1.06273
10	07	-1.3490352	1.8075806	5.53313	-11.34365	-2.46309
11	08	-1.3238584	1.7269712	5.46601	-18.21978	-5.12595
12	09	-1.2746349	1.5827714	5.34971	-25.92367	-8.38452
13	10	-1.2066696	1.4015248	5.16671	-30.73168	-10.90536
14	11	-1.1081768	1.1691900	4.79093	-30.30983	-11.77315
15	12	-0.9882935	0.9393943	4.09266	-18.73367	-7.46594
16	13	-0.8527599	0.6982302	3.51664	-13.58704	-5.79385
17	14	-0.6875841	0.4555011	2.87777	-8.12521	-3.45415
18	15	-0.5618981	0.2697318	2.65110	-13.46220	-9.19954
19	16	-0.3247234	0.0508937	2.05248	-9.73511	-10.91030
20	17	-0.3107299	0.1043721	1.72413	-1.64962	1.56380
21	18	-0.1772295	0.0129858	1.41982	-3.66060	-3.68489
22	19	-0.0653010	-0.0443728	1.37142	-3.40102	-9.72741
23	20	-0.0754223	-0.0154689	1.20130	-2.78914	-4.23148
24	21	0.0036170	-0.0071511	0.99791	-1.98240	-1.43283
25	22	0.1187272	-0.0102101	0.83572	-0.60822	-4.86125
26	23	0.2209112	0.0146637	0.72352	1.45842	-6.82762
27	24	0.2896986	0.0763799	0.51022	-0.54625	-1.50907
28	25	0.3704396	0.1191009	0.42920	1.42651	-3.62492
29	26	0.4859078	0.2080054	0.34326	4.43362	-5.62021
30	27	0.5252493	0.2734521	0.22598	-0.43797	-0.48694
31	28	0.6496628	0.4010884	0.16672	4.74957	-4.19467
32	29	0.6504449	0.4171563	0.12570	0.84173	-1.18445
33	30	0.7402801	0.5398832	0.07407	1.88838	-1.62629
34	31	0.8187324	0.6579482	0.04817	3.69005	-2.47491
35	32	0.8567829	0.7355403	0.02073	-0.78796	0.29268
36	33	0.8978652	0.8033260	0.01124	0.81423	-0.56718
37	34	0.9527803	0.8998400	0.00855	2.93553	-1.59007
38	35	0.9428425	0.8867987	0.00373	0.69778	-0.43066
39	36	0.9605985	0.9222061	0.00158	0.13000	-0.10868
40	37	0.9894981	0.9779504	0.00024	0.43655	-0.23120
41	38	0.9953827	0.9904516	0.00003	0.12419	-0.06702
42	39	0.9995176	0.9989702	0.00000	0.02510	-0.01304
43	40	0.9998959	0.9997996	0.00000	-0.00332	0.00156
44	41	1.0000323	1.0000627	0.00000	0.00077	-0.00035
45	Number of f calculations:			207		

46 **Number of gradient calculations:** 198

The result indicates that when `eta` is higher, the line search is less rigorous, so it requires more steps to reach the 2-D local minimum. However, when `eta` is higher, the number of function value and gradient calculations is lower since the less rigorous line search, too.